



14 ways to build RAG, one repo

Every major Retrieval-Augmented Generation pattern, each a runnable project on a shared core.



One shared core, 14 pluggable pipelines

Every pattern implements the same interface: `index(docs)` and `query(q)`
→ `answer + sources`.



Which RAG, when?

Start simple and measure. Add complexity only where the baseline fails.

Just getting started / need a benchmark	■ Naive
Exact terms: product codes, IDs, error codes	■ Hybrid
Vague or chatty user questions	■ Advanced · Multi-Query
Short questions, very different doc wording	■ HyDE
Chunks mean nothing on their own ("it", "the plan")	■ Contextual · Parent-Doc
Follow-up questions in a chat	■ Conversational
Hallucinations are expensive	■ Corrective · Self-RAG
Mixed simple + complex traffic	■ Adaptive
"Who designed the thing that X uses?"	■ Graph · Agentic
Users must verify every claim	■ Citation

Rule of thumb: **fix retrieval before touching the prompt or the model**. If the right chunk never reaches the LLM, no amount of prompting can save the answer.