



14 ways to build RAG, one repo

Every major Retrieval-Augmented Generation pattern, each a runnable project on a shared core.



One shared core, 14 pluggable pipelines

Every pattern implements the same interface: `index(docs)` and `query(q)`
→ `answer + sources`.



Which RAG, when?

Start simple and measure. Add complexity only where the baseline fails.

Just getting started / need a benchmark

■ **Naive**

Exact terms: product codes, IDs, error codes

■ **Hybrid**

Vague or chatty user questions

■ **Advanced · Multi-Query**

Short questions, very different doc wording

■ **HyDE**

Chunks mean nothing on their own ("it", "the plan")

■ **Contextual · Parent-Doc**

Follow-up questions in a chat

■ **Conversational**

Hallucinations are expensive

■ **Corrective · Self-RAG**

Mixed simple + complex traffic

■ **Adaptive**

"Who designed the thing that X uses?"

■ **Graph · Agentic**

Users must verify every claim

■ **Citation**

Rule of thumb: **fix retrieval before touching the prompt or the model**. If the right chunk never reaches the LLM, no amount of prompting can save the answer.